

AI Product Manager | Associate Product Manager | Product Analyst

Computer Science graduate and NextLeap Product Management Fellow who has built four AI products end to end, with expertise in LLMs and RAG. At Bosch, replaced manual reporting with self-serve dashboards adopted by 3 teams, cutting reporting effort ~85%. I ship AI features and define the metric that decides whether they launch.

Skills

Product Management: Product discovery, user research, PRDs, MVP scoping, RICE prioritization, roadmapping, North Star metrics, A/B experimentation

AI / Machine Learning: LLMs, RAG, prompt engineering, AI evaluation, LLM guardrails, embeddings, semantic search (FAISS), OpenAI API

Data & Analytics: SQL, Python (Pandas), Tableau, ETL automation, funnel & activation analysis, KPI design, REST APIs

AI Product Portfolio**FundBot** — Zero-hallucination assistant for regulated finance

GitHub

Python, RAG, FAISS, Sentence Transformers, OpenAI API

- Grounded every answer in **300+ Mirae Asset fact sheets** with source URL and date attached, so users verify claims against the document instead of trusting the model.
- Built a **41-query labeled evaluation set** as the release gate instead of eyeballing outputs: **92.7% top-1 retrieval accuracy, zero out-of-scope answers, 111 regression tests.**
- Added PII detection, advice-blocking, and scope validation with an extractive fallback, so a model failure returns a sourced fact, not a guess.

Recruiter Outreach Intelligence — Ranked, trackable job-search pipeline

GitHub

Python, FastAPI, PostgreSQL, Streamlit, LLM

- Reframed the goal from “apply to more roles” to “reach the right recruiter on a genuinely open role,” tracking every contact in a state machine.
- Built **4 scoring engines** (job relevance, hiring activity, recruiter relevance, outreach priority) so output is a **ranked queue** rather than a list to triage — with **257 offline tests.**

AI Shopping Assistant — Conversational grocery commerce

GitHub

Next.js, TypeScript, Firestore, LLM

- Converts intent (“dinner for four under Rs 500”) into cart-ready baskets, with **query-to-cart conversion** as the North Star measured against a keyword baseline.
- Constrained the model to the live catalog via schema validation and keyword fallback to **eliminate out-of-catalog suggestions** — the trust failure behind AI-search abandonment.

AI Discovery Engine — Voice-of-customer opportunity ranking

GitHub

Python, LLM, Pandas

- Clusters **800+ reviews from 6 sources** into ranked pain points traceable to source reviews, compressing multi-day research to about an hour (est.).

Experience**Bosch Automotive Electronics** — Data Analyst Intern

Bengaluru | Jun-Dec 2024

- Ran discovery with **8+ line engineers and quality leads** across three teams keeping separate manual KPI reports, then scoped 4 dashboard modules sequenced by decision impact rather than request volume.
- Shipped unified Tableau dashboards covering **15+ machines across 5 production lines, adopted by 3 teams within 3 weeks**, cutting recurring manual reporting ~85% (~10 hrs/week, est.) via self-serve access.
- Defined 4 reliability KPIs (MTTR, MTBF, Technical Availability, Technical Loss) on an automated Python ETL pipeline over Oracle SQL and HiveQL, letting engineers rank defect drivers and cutting root-cause investigation **from days to hours.**

Education & Certifications**Vellore Institute of Technology, Bhopal** — B.Tech, Computer Science & Engineering

2021-2025 | GPA 7.7/10

NextLeap — Product Management Fellowship, Cohort 47

Mar-Jun 2026

Self-directed teardowns and growth PRDs: Eureka by ALLEN, voice input in urban India, n8n onboarding.

Certifications: PM Simulation (EA, Forage) | Quantitative Research (JPMorgan, Forage) | ML Specialization (Coursera)